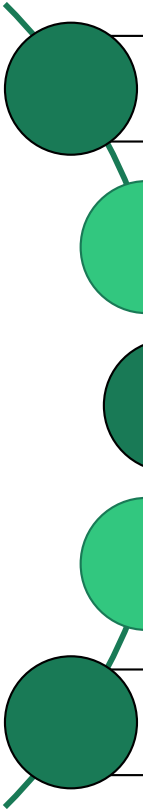




# Spotify Recommendation Algorithm

# Our Agenda For Today



How We Came Up With Our Idea and Found Our Dataset

Different Approaches to Recommendation Algorithms

Our Spotify Recommendation Algorithm: What It Does and How It Works

How Does Our Algorithm Perform?

Additional Topics for Further Discussion and Improvement



# How We Came Up With Our Idea and Found Our Dataset

# What Brings Us Together?

## Potential Datasets

- › Art and Artists
- › Netflix Shows
- › Open Data on AWS
  - › Phishing
- › Spotify Hit Predictor



# Our Dataset – 169,909 Songs from 1921-2020

The datasets features are:

<u>Acousticness</u>	<u>Artists</u>	<u>Danceability</u>	<u>Duration</u>
<u>Energy</u>	<u>Explicit</u>	ID	
<u>Instrumentalness</u>	<u>Key</u>	<u>Liveness</u>	<u>Loudness</u>
<u>Mode</u>	<u>Name</u>	<u>Popularity</u>	
Release date	<u>Speechiness</u>	<u>Tempo</u>	<u>Valence</u>
	<u>Year</u>		



# Different Approaches to Recommendation Algorithms

## We Decided Against Our Initial Idea of Collaborative Filtering Due to a Lack of User Data

- Collaborative Filtering through Matrix Factorization is a common aspect of services like Netflix or Spotify.
  - It relies on similarities in taste and between songs to determine possible likes/matches.
- We didn't use this, since this requires user data on song preference, which was not something we could get.


Preference	Song A	Song B	Song C	Song D
User A	1	2	1	1
User B		1	1	3
User C	1		3	2




# We Chose to Use Similarity Metrics for Our Recommendations

- To find song similarity, we needed to use metrics which calculated the distance between the numerical features of our data.
- Manhattan distance was our first choice, since it works better in higher dimensions, and is less sensitive to extreme values
- Cosine similarity, while not equivalent to Euclidean distance, gives us something similar.
  - More on that later...

Manhattan Distance



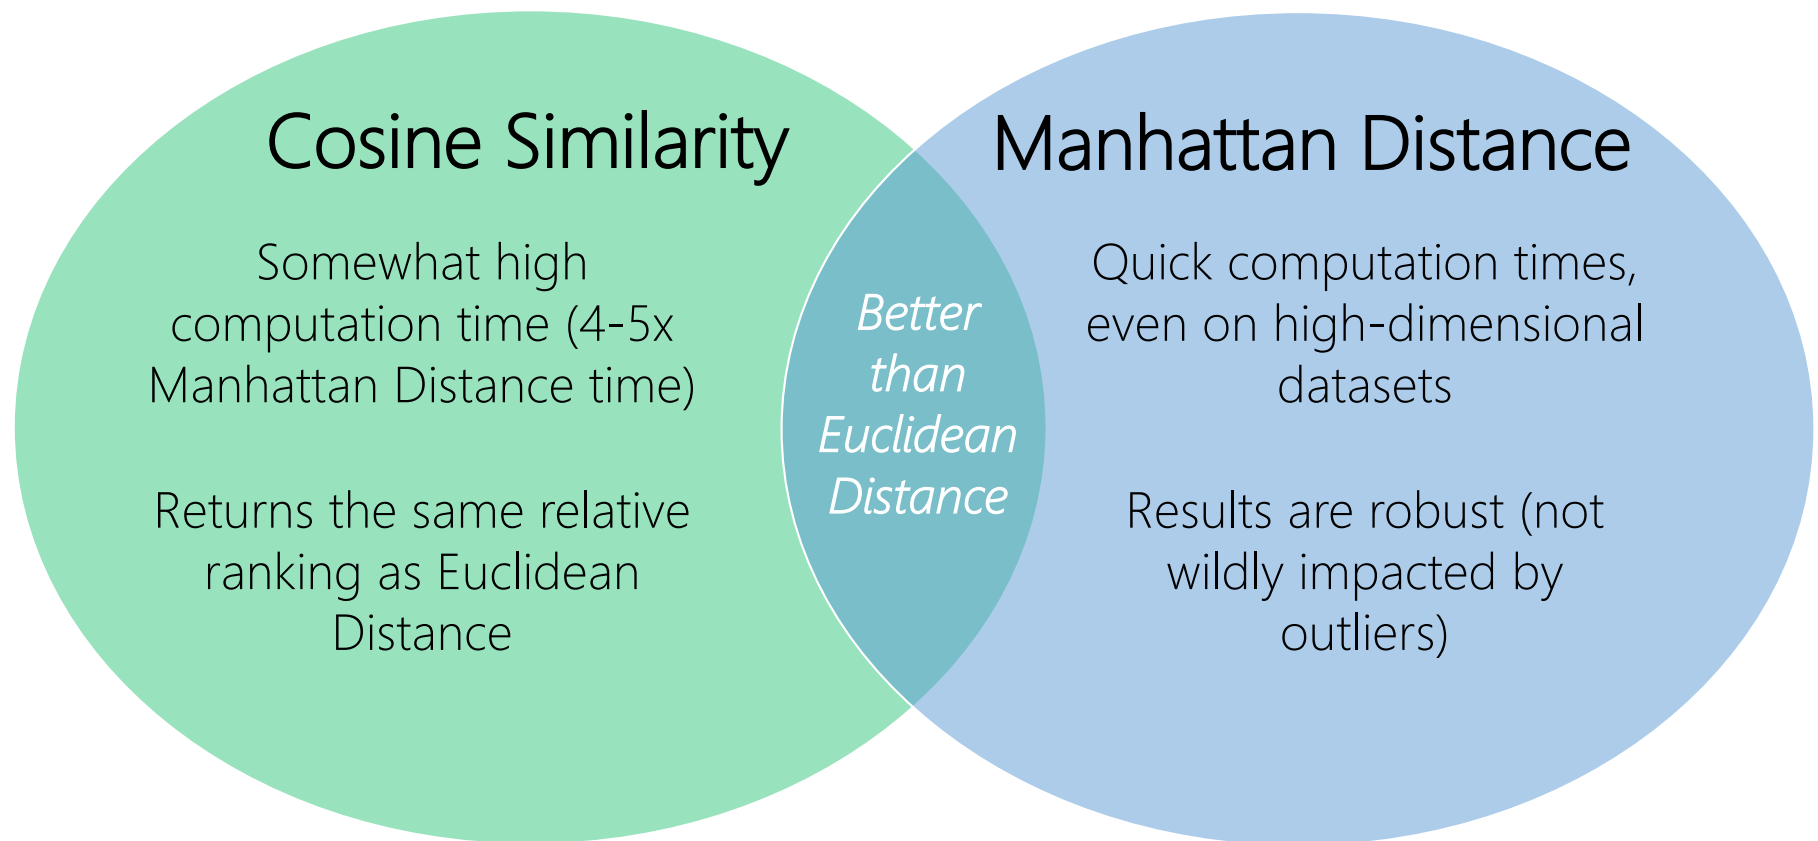
Cosine Similarity





# Our Recommendation Algorithm: What It Does and How It Works

## We Built Two Different Recommendation Functions



# Our Data Preprocessing Steps



**Drop Unnecessary  
Columns**

**Drop Duplicate  
Data Entries**


**Convert Categorical  
Variables**

**Min-Max Scaling  
to Normalize**




**End Result: a Clean Dataset**

# Dealing with High Correlations Between Features



## After Dropping the Energy Column



# Exploring Our Data: KPrototypes Clustering

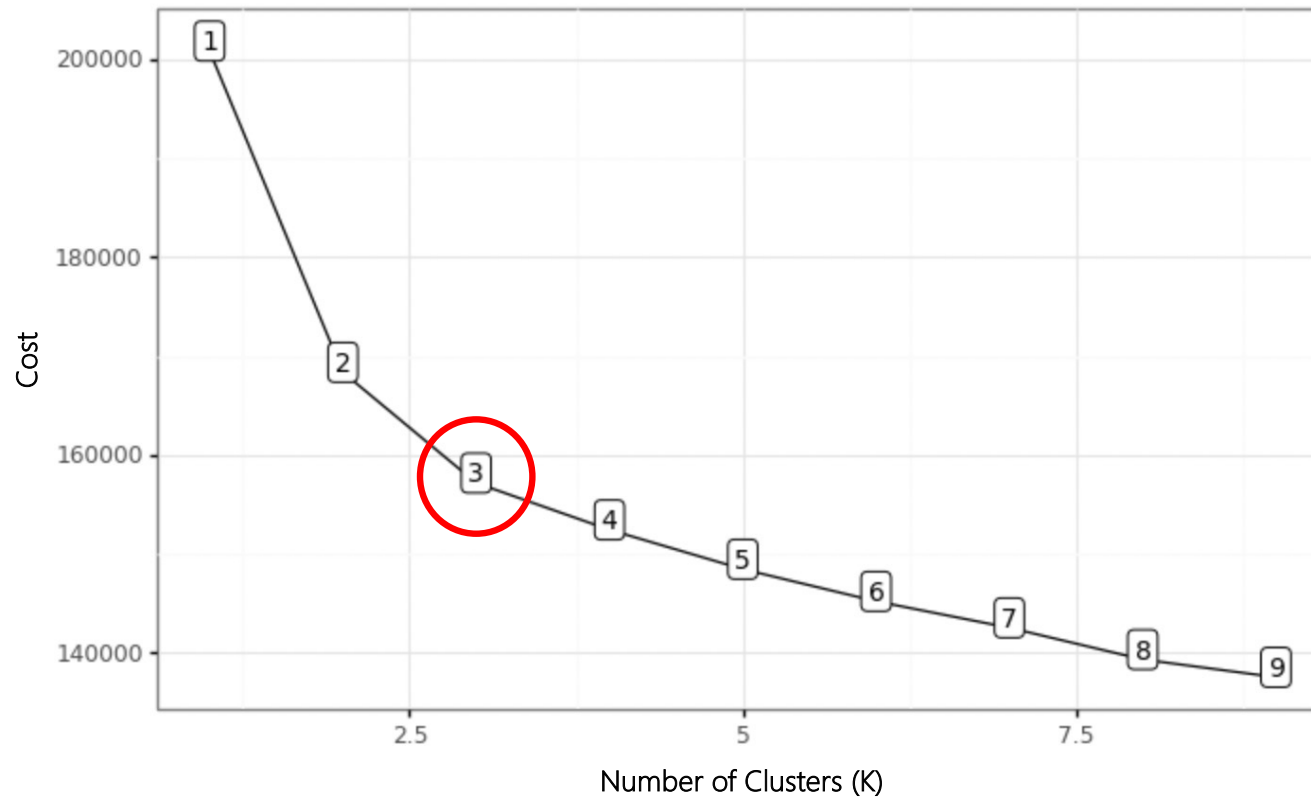
We chose to use KPrototypes clustering for our data to improve the computation speed and accuracy of the song recommendation functions.

	Accepts Numerical Data	Accepts Categorical Data	Time Complexity	Interpretability
KMeans	✓	✗	Decent	Very Good
KModes	✗	✓	Very Good	Unusable
<b>KPrototypes</b>	✓	✓	<b>Awful</b>	Decent



# Choosing the Correct Number of Clusters

Optimal Number of Clusters Using Elbow Method



# Visualizing Our Clusters Using PCA

0

**Cluster 0:** classical music, symphonies, podcasts, radio, and white noise


1

**Cluster 1:** music those using our recommender are likely to listen to

2

**Cluster 2:** older music and songs in non-English languages

Cluster Visualization in Two Dimensions





How Does Our  
Algorithm Perform?

# You Be The Judge! Song Recommendations for Classmates

<u>A</u>	<u>B</u>	<u>C</u>	<u>D</u>
Until You Come Back to Me – Aretha Franklin  Feel Like I Do – Disclosure  I Like It – DeBarge  Ultimate – Denzel Curry  Blue World – Mac Miller	Streetcar – Daniel Caesar  Nights – Frank Ocean  Die for You – The Weeknd  Open – Kehlani  Garden – SZA	More – Chief Keef  Love Sosa – Chief Keef  15 – Taylor Swift  Going Bad – Drake  Poundcake – Drake/Jay-Z	Sunflower – Post Malone  Beautiful People – Ed Sheeran  I Don't Care – Ed Sheeran  As Long As You Love Me – Justin Bieber  Scared of the Dark – Lil Wayne
<u>Manhattan</u> Right Now – Lil Uzi Vert Yours Truly – Post Malone  <u>Cosine</u> For However Long – Bryson Tiller Without Me – Halsey, ILLENIUM	<u>Manhattan</u> Love Don't Change – Jeremih To Die For – Kygo, St Lundi  <u>Cosine</u> Hurting – Kygo The Worst – Jhene Aiko	<u>Manhattan</u> Cocoa Butter Kisses – Chance the Rapper Broken Clocks – SZA  <u>Cosine</u> Clique – Kanye West 3500 – Travis Scott	<u>Manhattan</u> Graveyard – Halsey Like I Can – Sam Smith  <u>Cosine</u> Reborn – Kids See Ghosts Canyon Moon - Harry Styles



# Get Your Own Custom Song Recommendations!

Using Google API, we can import data from our Google Sheet to our Colab File and automate the recommendation system!




# Measuring Accuracy Quantitatively is Difficult Without True Labels

- Common error metrics such as MAE, RMSE, MSE require there to be a **predicted label** and an **actual label**.
  - Because our dataset only contains information regarding the songs/artists/albums, we were unable to product any meaningful quantitative accuracy measure.
- However, given more resources, there are some steps we could take to create an accuracy system:

## 1. Internal Validation

- Cohesion across samples
- Compute the similarity between several records via 'Euclidian distance' in featurespace



$$\text{Cohesion}(C_k) = \sum_{x \in C_k; y \in C_k} \text{similarity}(x, y)$$

## 2. External Validation / Twin Sample Validation

- Creating a sample of records that is expected to exhibit similar behavior to the training data
- Perform unsupervised learning on the 'twin-sample'
- Compare similarity between two sets results
  - F1, Jaccard Similarity.... any kind of External Validation method



# Additional Topics for Further Discussion and Improvement

# Possible Additions and Improvements to Our Project

## Up-to-Date Song Dataset

1

Our current dataset stops at 2020. We looked into using Spotify's API to pull up-to-date song data, but found it was more trouble than it was worth.

## Web-Based Front-End

2

Would allow users to input their songs and receive real-time results through a sleek user interface. Unfortunately, we couldn't get it working in time.

## Artist Recommender

3

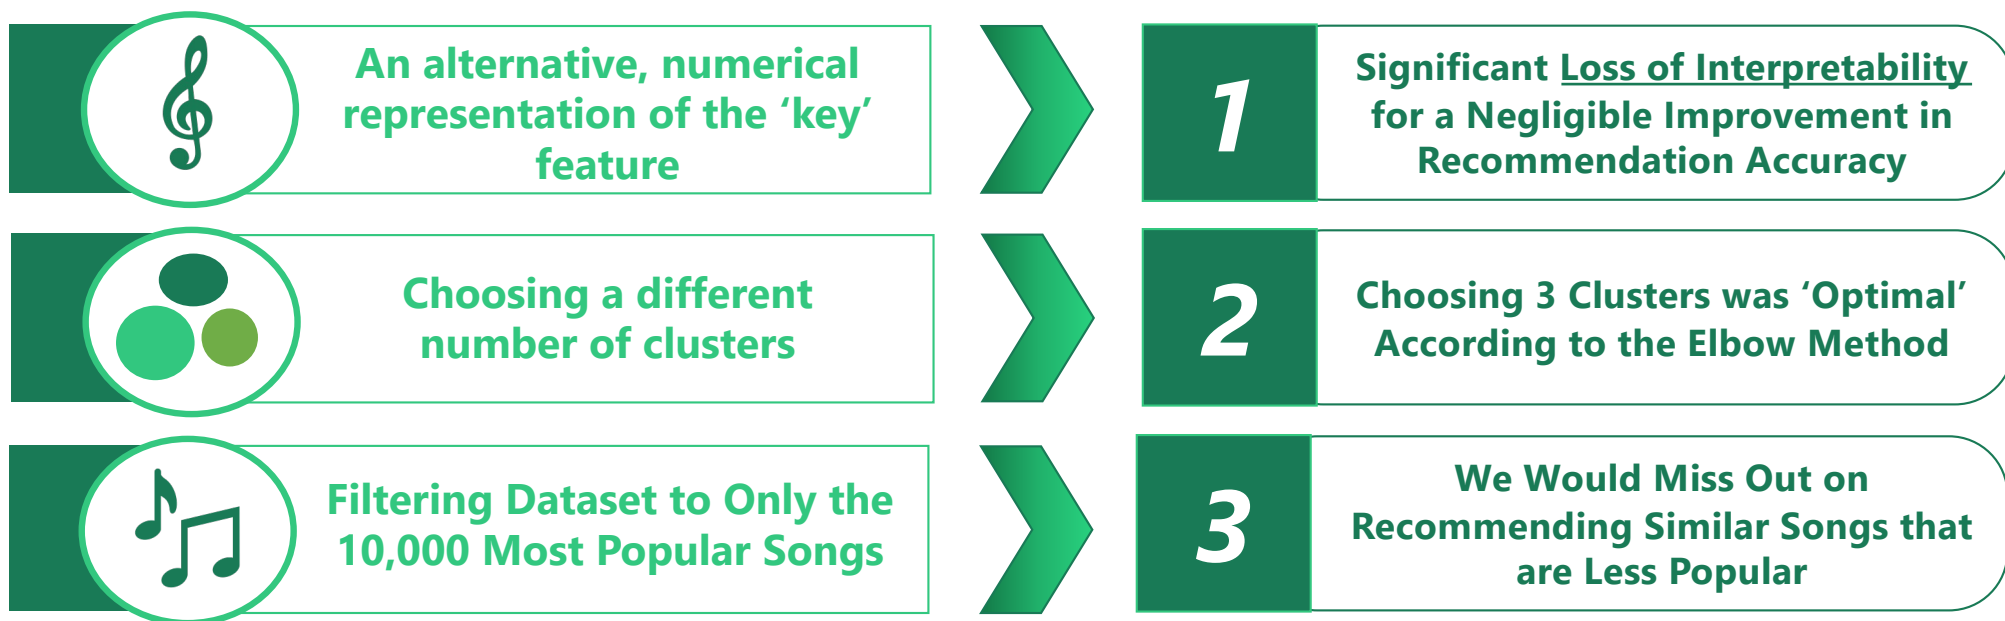
Similar to the song recommender. It would use naïve weighting for each artist and artist-relative popularity for each artists' songs.



# Possible Ways to Improve the Underlying Algorithm

We feel our algorithm performs **as best as can reasonably expected** in most general-case scenarios, considering we used unsupervised learning to generate recommendations from imperfect data.

A few suggestions for potential improvements we could make to the underlying algorithm are:



Thank You!

Any Questions?